


Colossal-AI accelerates Protein Structure Prediction by up to 11 times



● Experimental result
● Prediction

The Colossal-AI team and BioMap have created a new tool called xTrimo Multimer that can quickly predict the 3D shape of both single proteins (monomers) and multiple proteins that work together (multimers) using AlphaFold but optimized for a faster runtime.

**Reduced training time
from 11 days to only 67 hours**

**Accelerated inference time
by up to ~11.6 times**

“The collaboration with HPC-AI Tech brings together the cutting-edge technology in large AI model training from Colossal-AI and the biocomputing domain expertise from BioMap. The release of xTrimo Multimer is an important step towards integrating the advantages of large AI model training and inference into the construction of BioMap’s xTrimo multimodal system.”



Le Song

Chief AI Scientist, BioMap

About the customer

BioMap is a team of world-renowned scientists who have extensive expertise in disease biology, bioinformatics, machine learning/deep learning, and antibody engineering. BioMap was co-founded by Baidu's Founder/CEO Robin Li and the former CEO of Baidu Ventures, Wei Liu. They are committed to bringing first-in-class medicine for unmet medical needs in the areas of immune-oncology, autoimmune diseases, fibrosis and aging-related diseases.



Industry: Biotechnology



Solution: Performance optimization for deep learning to predict protein structures



Model: AlphaFold



Product: Colossal-AI FastFold

The challenge

Protein structure prediction is one of the most important topics in structural biology and supplements the understanding of gene translation and protein function. It is a process where scientists try to figure out the 3-dimensional shape of a protein from its genetic code (amino acid sequence). Knowing the structure of a protein is important because it can reveal how the protein works and how it interacts with other molecules. This information can be used to design new drugs, understand and treat diseases, and develop new technologies.

Unfortunately, the multi-level structure and sophisticated protein interactions make it challenging to predict the 3D structure accurately. In recent years, the success of deep neural networks has transformed various practices. Since


DeepMind's release of AlphaFold (which accurately predicts protein structure based on amino acid sequences), the field of biology has witnessed a boom in the utilization of AI for protein structure prediction.

Specifically, AlphaFold can generate end-to-end 3D structure predictions of protein monomers directly from amino acid sequences. The use of AlphaFold also exceeds the realm of monomers. Since the majority of proteins function as multimers, DeepMind's AlphaFold-Multimer model is recently released to be able to predict the structure of multimers.

The solution

To boost the development of AlphaFold, the Colossal-AI team has already released FastFold, an open-source and optimized implementation of AlphaFold. The Colossal-AI team managed to successfully minimize AlphaFold's training time from 11 days to only 67 hours and accelerate inference time by up to ~11.6 times. The Colossal-AI team continues its efforts to democratize large-scale AI model applications in the pharmaceutical field.

Interactions between proteins are critical to their biological functions. To address difficulties related to protein monomer and multimer structure prediction, the Colossal-AI Team proposed the industry's latest solution, the xTrimo Multimer. xTrimo Multimer is able to better reflect protein interactions, thus enhancing the potential target analysis, protein structure prediction/simulation, as well as high-precision antibody designs in drug discovery and development.

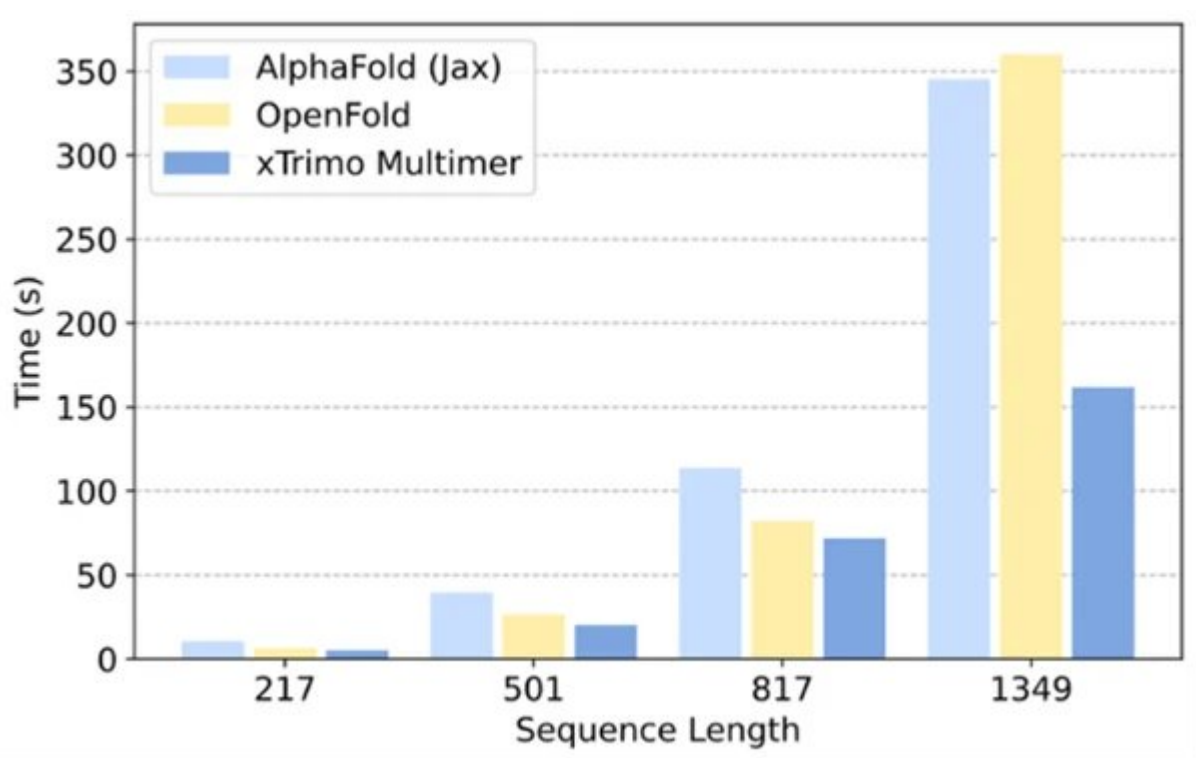


Prediction Result of AlphaFold2-Multimer and xTrimo-Multimer (Protein 3G6D)

The unaffordable economic and time costs from AlphaFold's inference has led to challenges in its research and development, particularly when facing long sequence inferences with rising computational complexity and memory consumption. Based upon computational features in the AlphaFold-Multimer model, the Colossal-AI team introduced CUDA optimization and Kernel Fusion techniques for the xTrimo Multimer achieving remarkable inference performances.

The results

Compared to AlphaFold2 and OpenFold (from Columbia University), the xTrimo Multimer has a significantly improved inference performance on a single GPU by 1.58-2.14 times and 1.14-2.23 times, respectively.



Additionally, the xTrimo Multimer model supports distributed inferences for lengthy sequences. After introducing Dynamic Axial Parallelism, the xTrimo Multimer was able to efficiently distribute computation and partial GPU memory across a variety of devices, thereby solving computational and memory challenges that long sequences face. xTrimo Multimer achieves a 8.47x and 11.15x speedup compared to OpenFold and AlphaFold 2 on multiple GPUs

with sequence lengths ranging anywhere from 2-3K. xTrimo Multimer also shows a 4.45x acceleration compared to Uni-Fold 2.0. Furthermore, xTrimo Multimer can support inferences with sequences reaching up to 4K, whereas OpenFold and AlphaFold 2 are restricted from such lengths due to GPU memory. With xTrimo Multimer, scientists can run a 4K length sequence inference in about 20 minutes.

Inference Latency for Extremely Long Sequence (s).

Sequence Length	AlphaFold	OpenFold	Uni-Fold	xTrimo Multimer (4 GPU)	xTrimo Multimer (8 GPU)
1816	840.2	638.44	335.79	102.82	75.37
3072	3328.9	3078.39	1829.05	710.154	431.17
4096	OOM	OOM	OOM	1919.66	1330.13

The future

"Our latest solution of protein Monomer and Multimer structure prediction is an important progress of Colossal-AI to solve the industrial problems in the real world. In the future, we will continue to cooperate with BioMap more deeply in biocomputing large models, to stimulate the application and implementation of deep learning in innovative drug development. "



Yang You

Founder & Chairman, HPC-AI Tech

Related content

Blog

Helping Drug Research and development — accelerating AlphaFold inference and training

Use a Laptop to Analyze 90% of Proteins, With a Single-GPU Inference Sequence Exceeding 10,000! Accelerate AlphaFold Inference by 5 Times and Reduce GPU Memory by 75%

Code

FastFold open source project on GitHub

Boost your AI business

Talk to an expert

Try open source

HPC-AI Tech offers vendor-supported, enterprise-grade solutions for Colossal-AI users, both on-premises and in the cloud. Our industry-leading technology accelerates and expands deep learning capabilities through parallel and

distributed training, inference, and fine-tuning of large neural networks using high-performance computing to reduce your costs. We operate globally and serve clients in various industries.



Newsletter

Sign up to get updates.

Resources

[Blog](#)

[Docs](#)

[GitHub](#)

[Slack](#)

[Forum](#)

[WeChat](#)

Copyright © 2023, HPC-AI Technology Inc.

All product and company names, as well as logos and trademarks or registered trademarks, are the property of their respective holders. Use of them does not imply any affiliation with or endorsement by them.